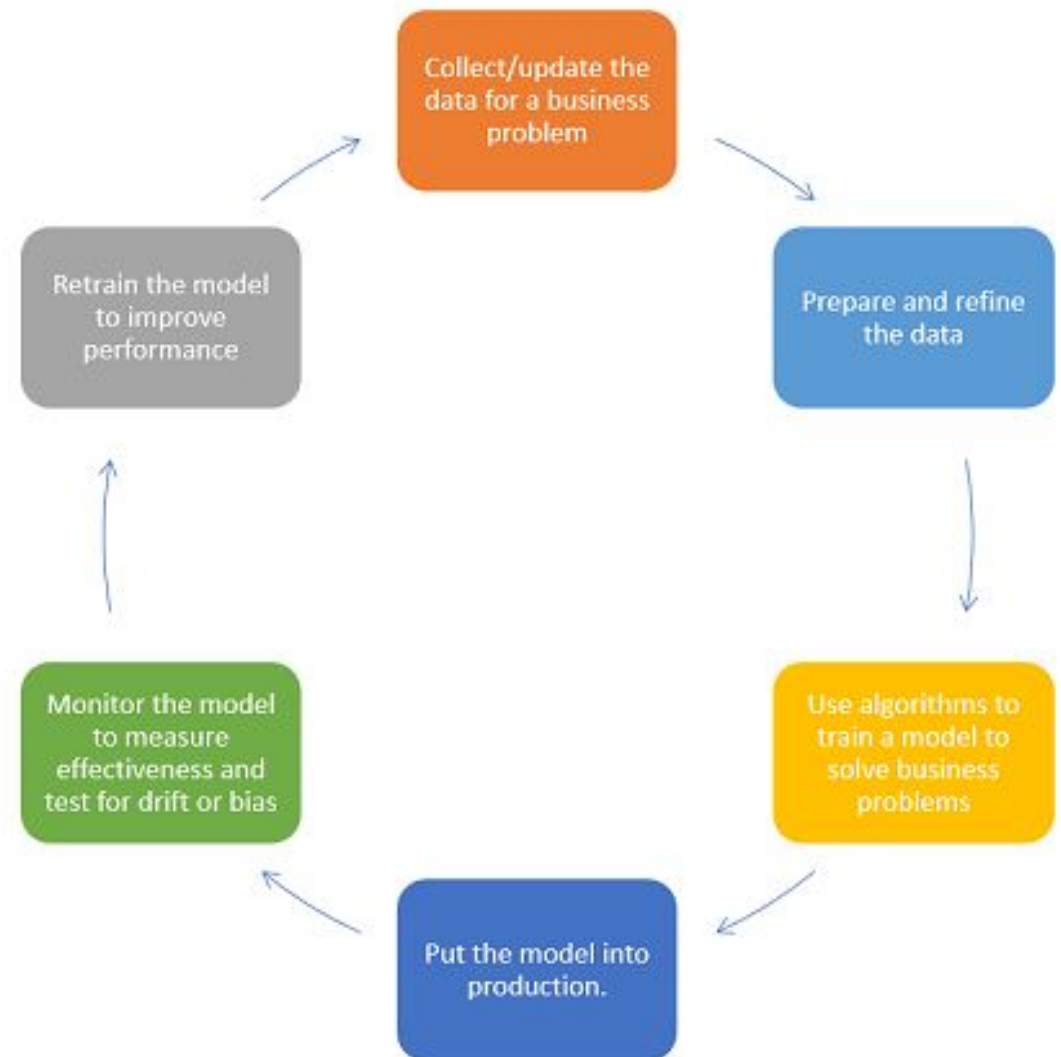


# ML Lifecycle

# Why understand ML Lifecycle and challenges encountered during the lifecycle?

## Most Enterprises Are Failing to Scale AI

In 2018, a team at Gartner research asked large enterprises the amount of AI adoption that they expect to deploy in the next 12 months. The enterprises said that they were planning to deploy AI in 23% of their systems. Then, in 2019, they went back to check how many of these projects actually deployed. The team found out that only 5% of the AI adoptions that enterprises wanted to deploy were actually deployed.



# 4 Phases to ML Lifecycle

**Phase 1** is Project Planning and Project Setup: At this phase, we want to decide the problem to work on, determine the requirements and goals, as well as figure out how to allocate resources properly.

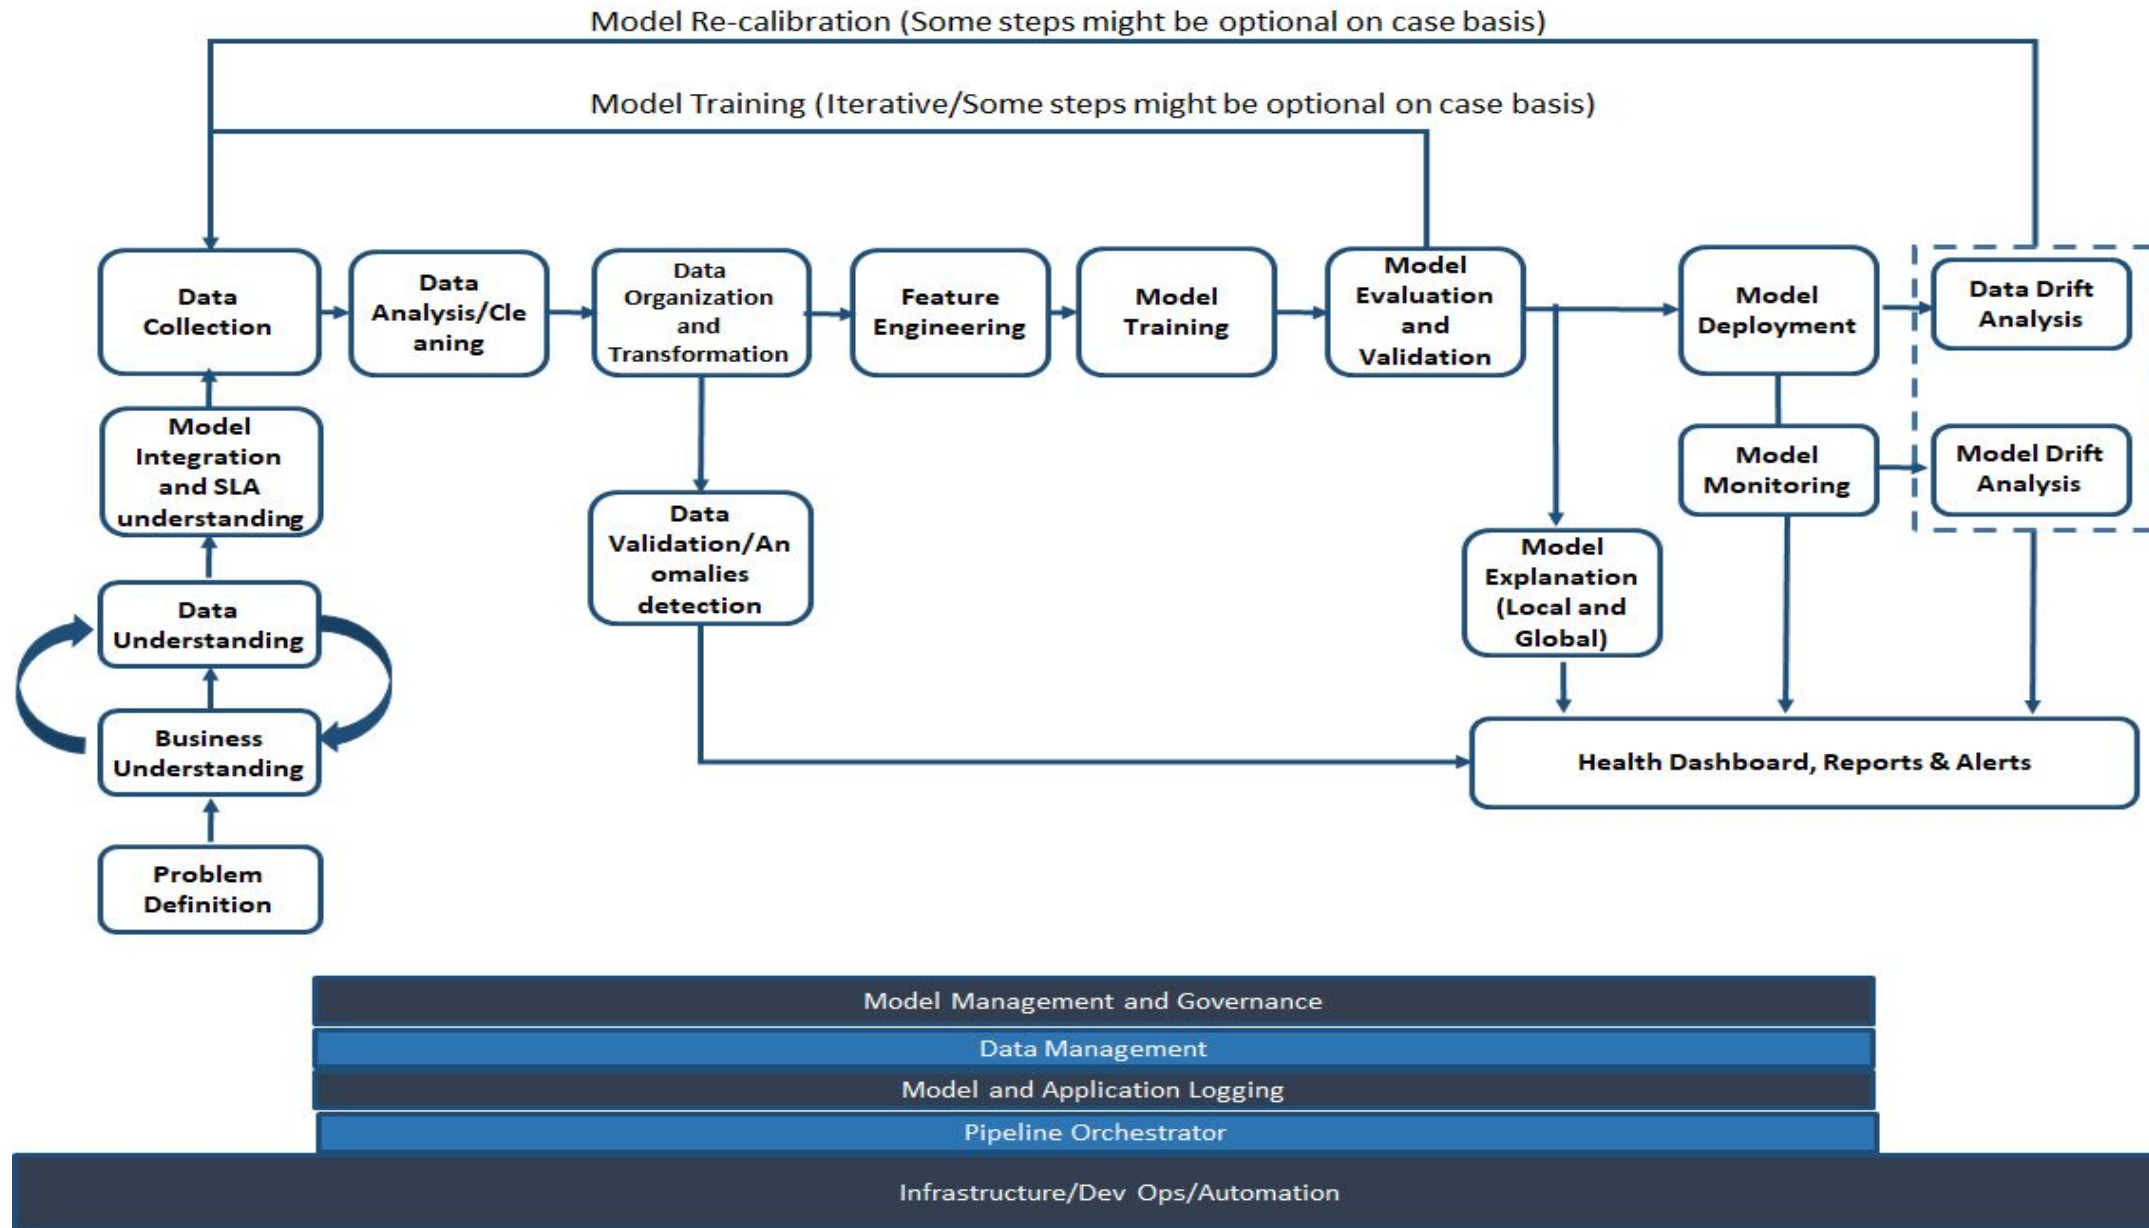
**Phase 2** is Data Collection and Data Labeling: At this phase, we want to collect training data (images, text, tabular, etc.) and potentially annotate them with ground truth, depending on the specific sources where they come from.

**Phase 3** is Model Training and Model Debugging: At this phase, we want to implement baseline models quickly, find and reproduce state-of-the-art methods for the problem domain, debug our implementation, and improve the model performance for specific tasks.

**Phase 4** is Model Deployment and Model Testing: At this phase, we want to pilot the model in a constrained environment, write tests to prevent regressions, and roll the model into production.

**Discussion: How does the real life ML lifecycle looks like?**

# Real world ML Lifecycle



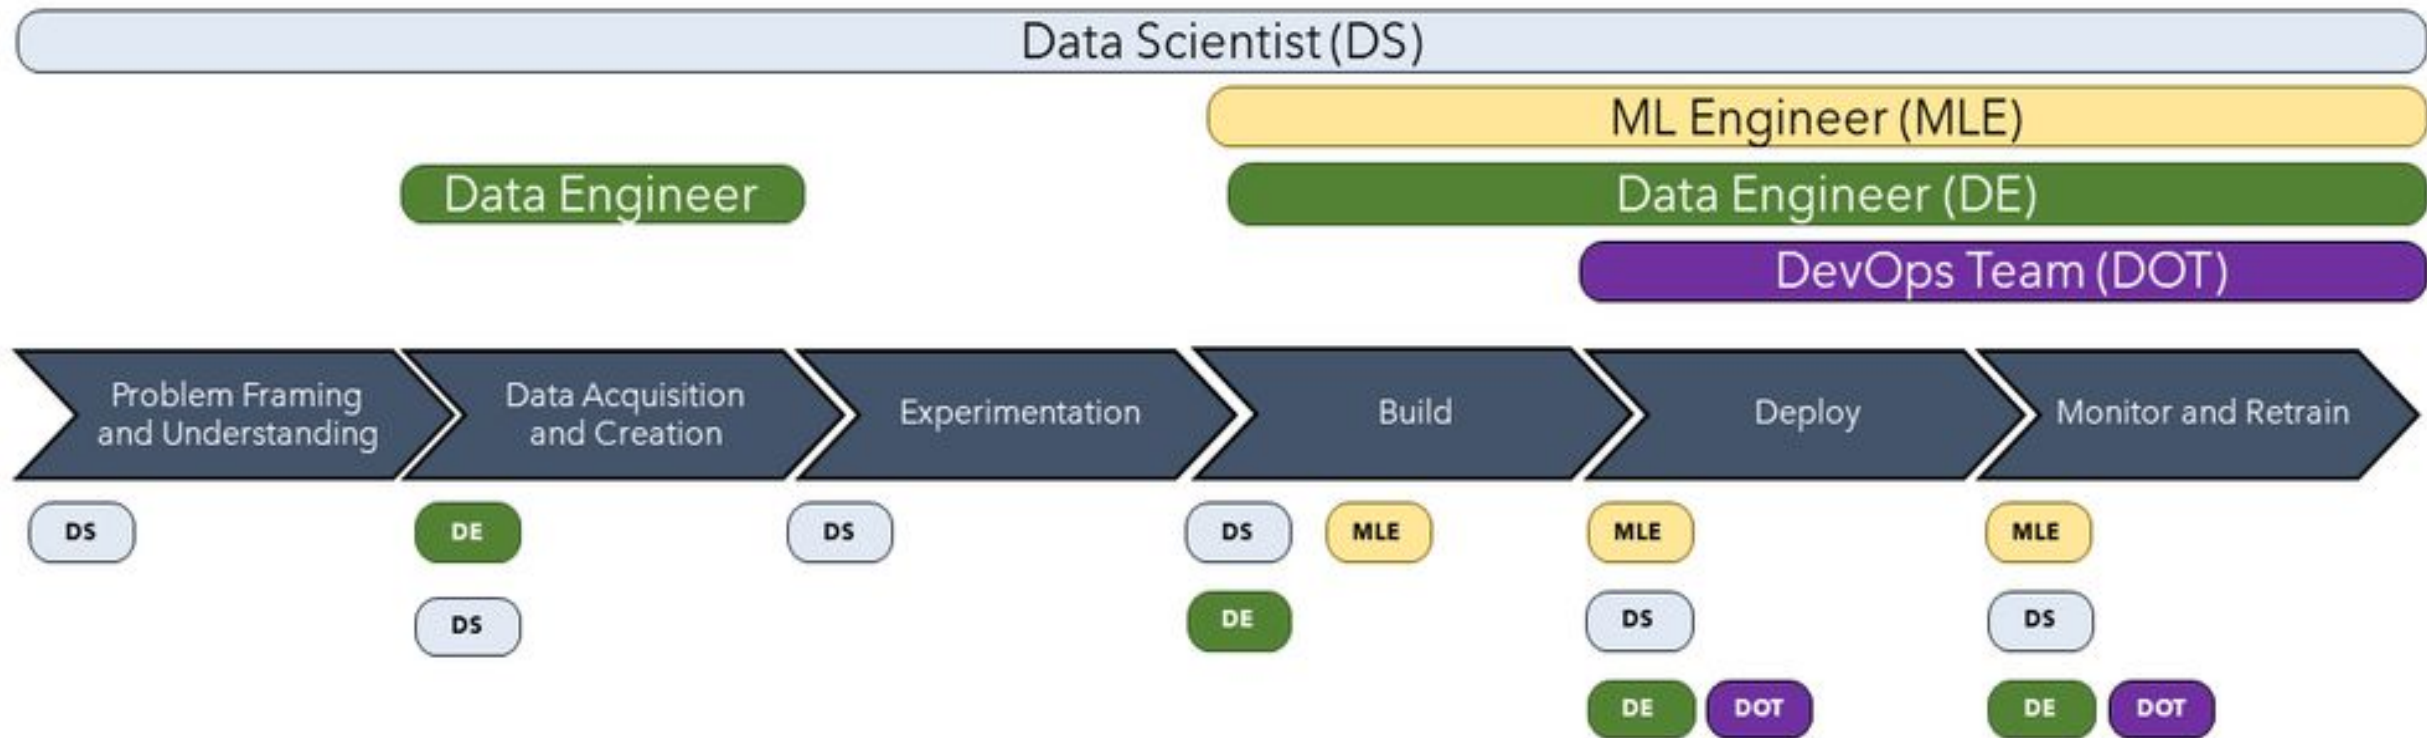
# Challenges with ML during development

- Development, training and deployment environment can be different
- Tools, libraries and dependencies can complicate deployment
- Tracking and analyzing experiment can become tedious to handle
- Difficult to reproduce experiment as input data changes
- ML Code end up in a spaghetti jungle

# Challenges with ML in production

- Live data is not equal to training data
- Feature engineering pipeline must match between training and serving infrastructure
- Seamlessly scale up and scale down deployed model
- Continuous training and champion challenger model deployment
- Different technology landscape between development and deployment

# Different Roles in ML Lifecycle





## References

- [An Introduction to MLOps: AI Engineering](#)

# Machine Learning Engineering

# Machine Learning Engineering (MLE)

## What is machine learning engineering?

Machine learning engineering is the process of using software engineering principles, and analytical and data science knowledge, and combining both of those in order to take an ML model that's created and making it available for use by the product or the consumers.

For example, a YouTube ML engineer might be in charge of developing the next generation YouTube recommendation algorithm and then developing an ML pipeline around it and integrating it into

Data Scientist

Machine Learning Engineer

Data Engineer



Research ML/AI  
Adv. Analytics



Operationalizing ML  
Optimizing ML



Adv. Programming  
Distributed Sys.

BDI

BIG DATA INSTITUTE

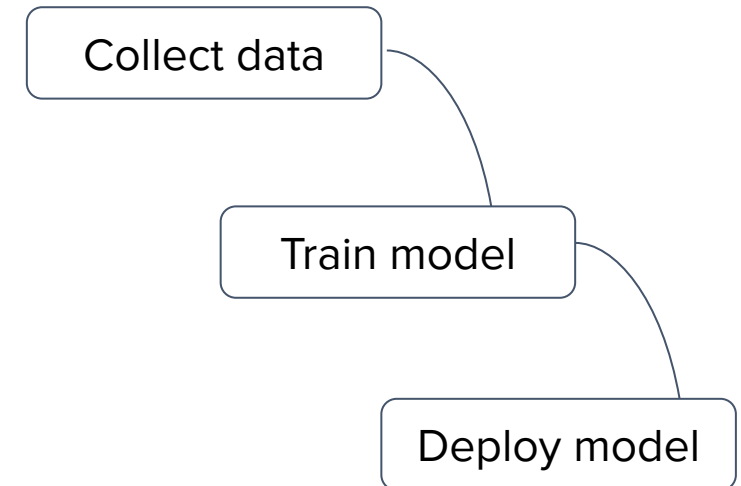
For more information go to <http://bigdatainstitute.io>

Attribution-NoDerivatives 4.0 International (CC BY-ND 4.0) Licensed



# ML in production: expectation

1. Collect data
2. Train model
3. Deploy model



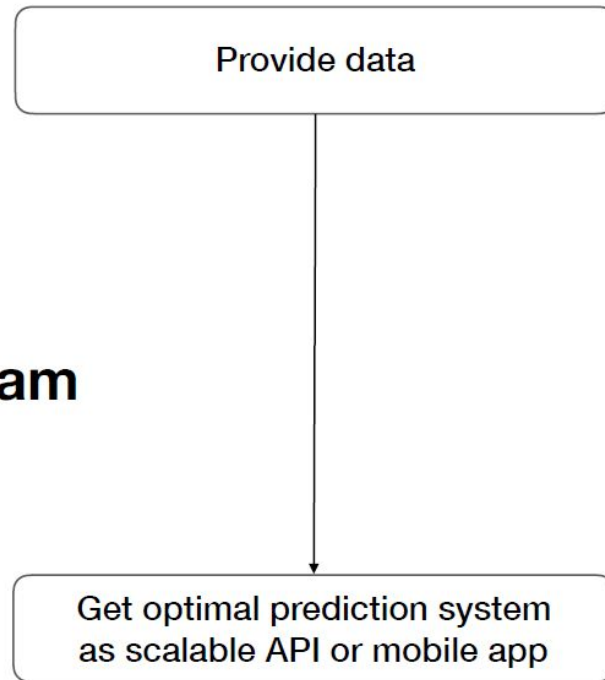
**Waterfall model**

# Problems?

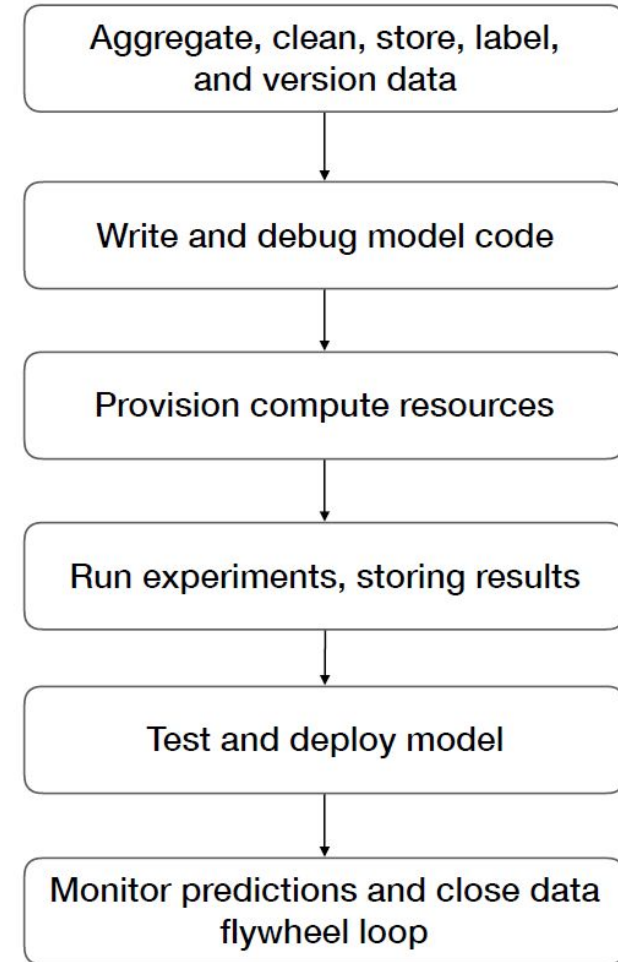
- What are the issues in ML in Production?
- Post on <https://padlet.com/rajaty6/bpmht0n2n8ls0d6>

# ML in Production

## Dream

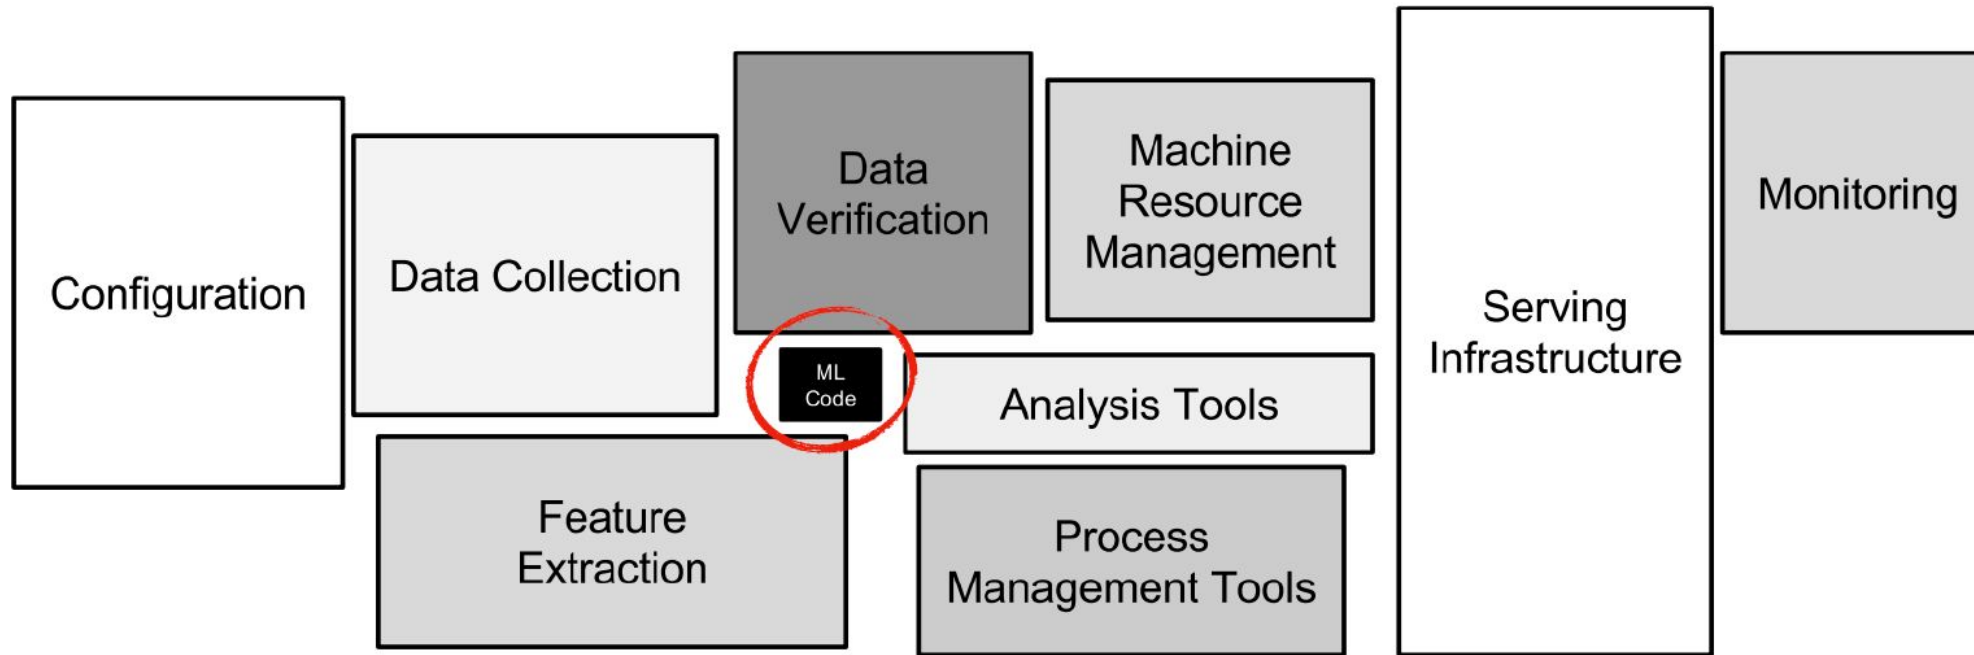


## Reality





# Machine Learning: Technical Debt



---

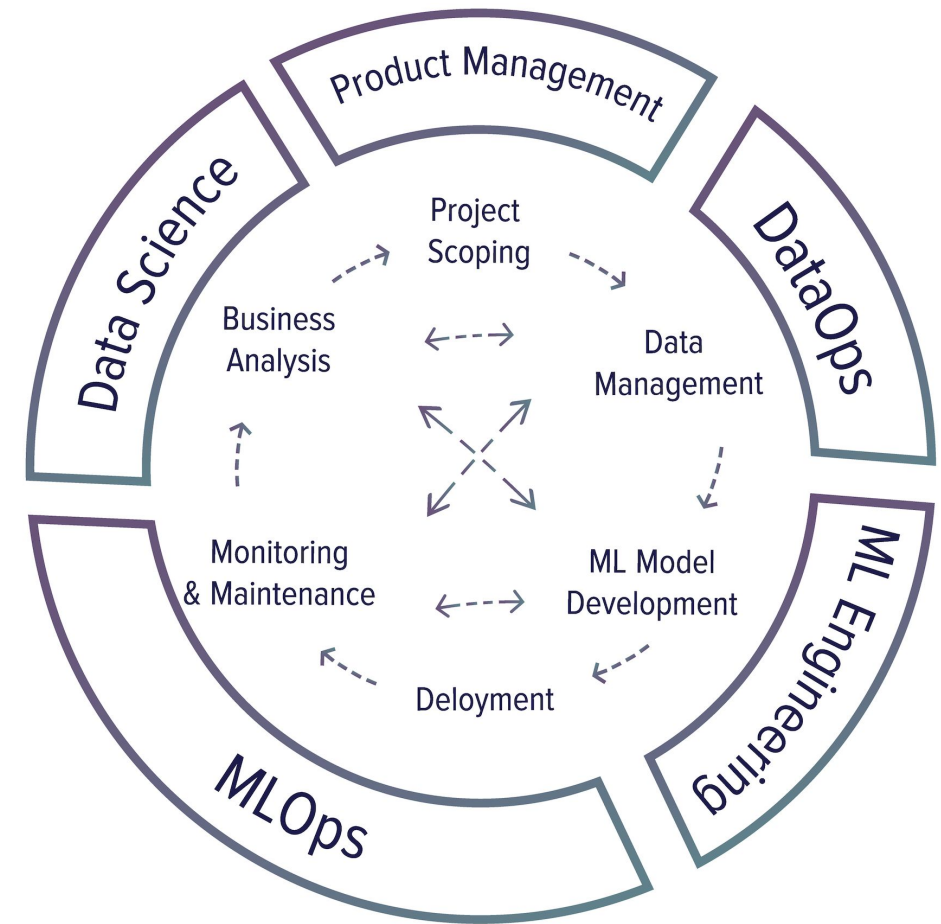
## Machine Learning: The High-Interest Credit Card of Technical Debt

---

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov,  
Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young

# ML in production: reality

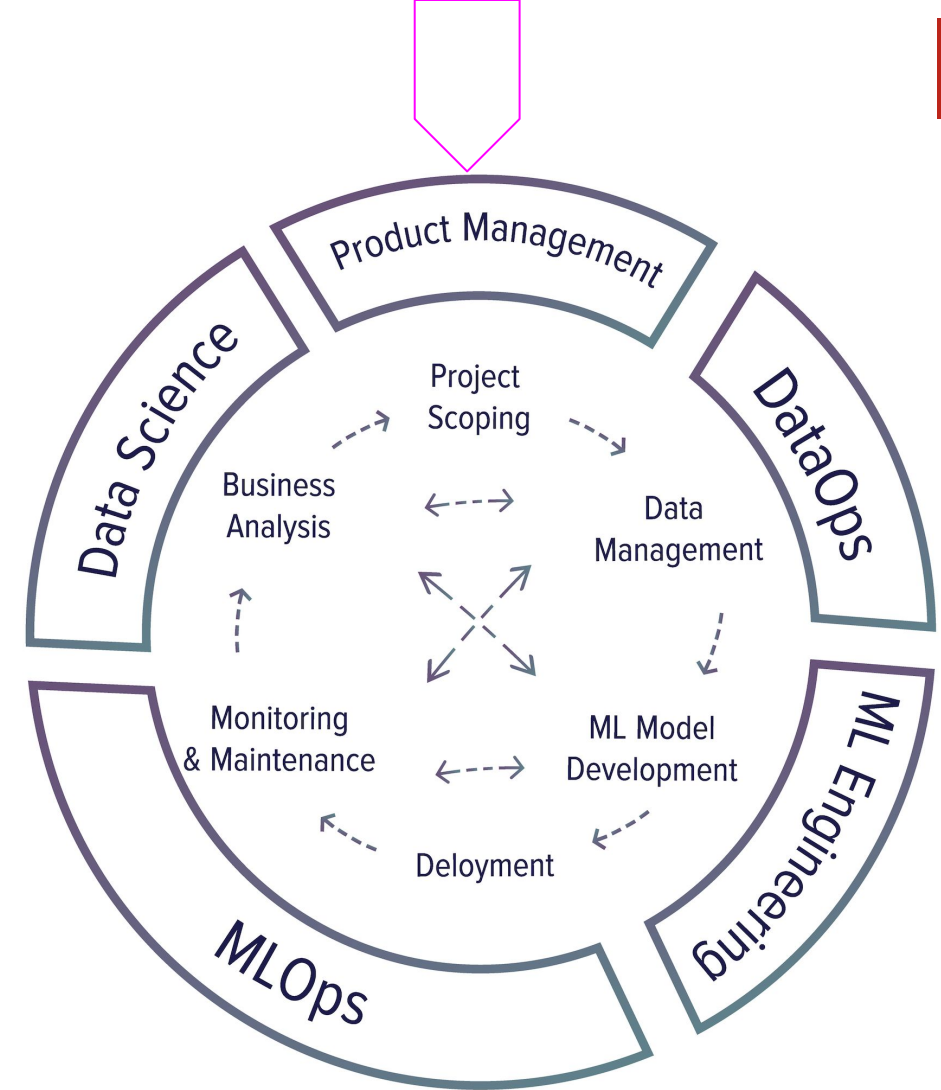
1. Choose a metric to optimize
2. Collect data
3. Train model
4. Realize many labels are wrong -> relabel data
5. Train model
6. Model performs poorly on one class -> collect more data for that
7. Train model
8. Model performs poorly on most recent data -> collect more recent data
9. Train model
10. Deploy model
11. Dream about \$\$\$
12. Wake up at 2am to complaints that model biases against one group
13. Get more data, train more, do more testing
14. Deploy model
15. Pray
16. Model performs well but revenue decreasing
17. Cry
18. Choose a different metric
19. Start over



**Iterative development**

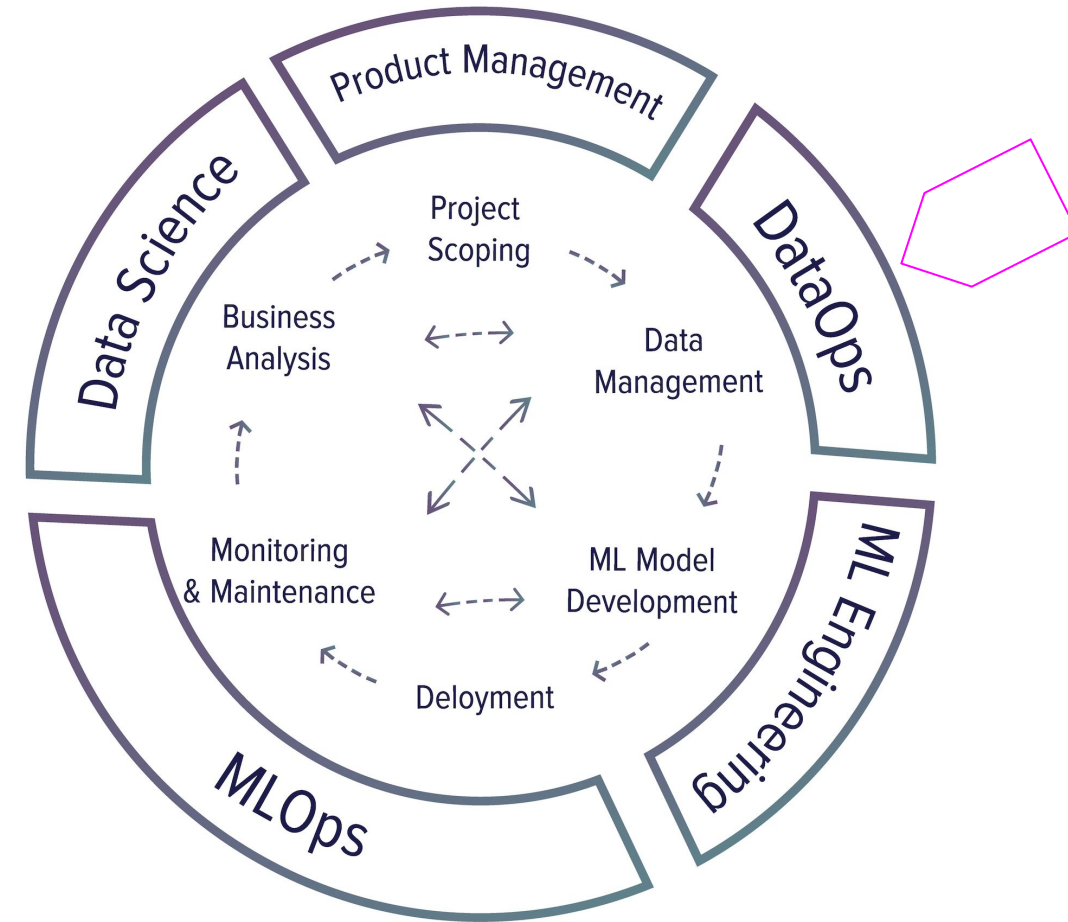
# Project scoping

- Goals & objectives
- Constraints
- Evaluation



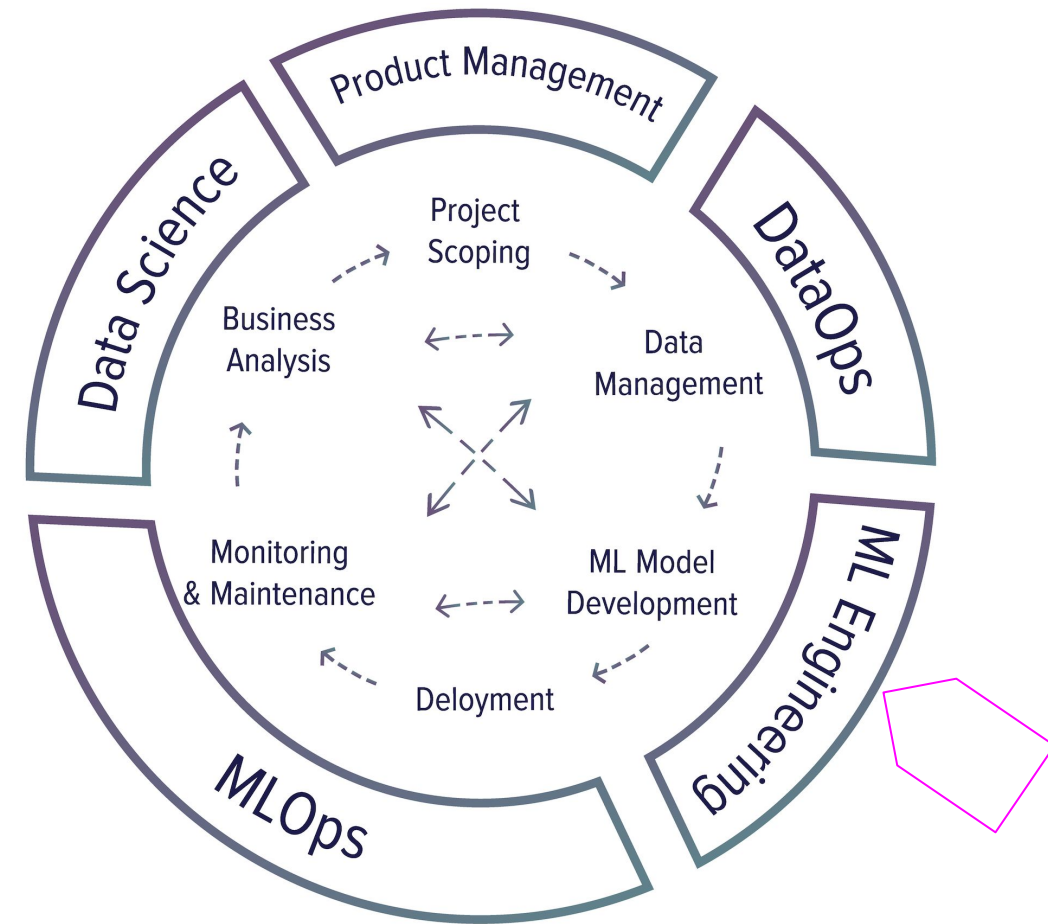
# Data management

- Data sources
- Data format
- Processing
- Storage
- Data consumer
- Data controller



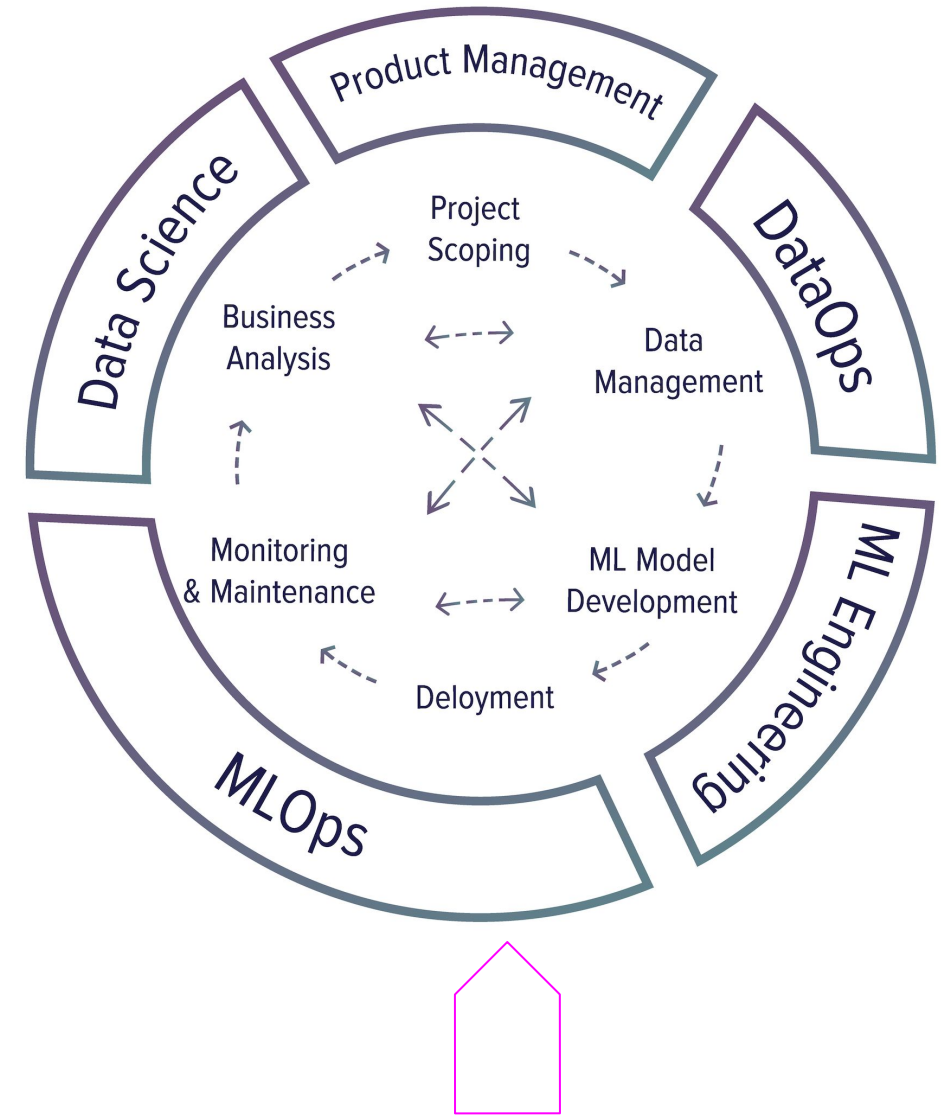
# Model development

- Dataset creation
- Feature engineering
- Model training
- Offline model evaluation



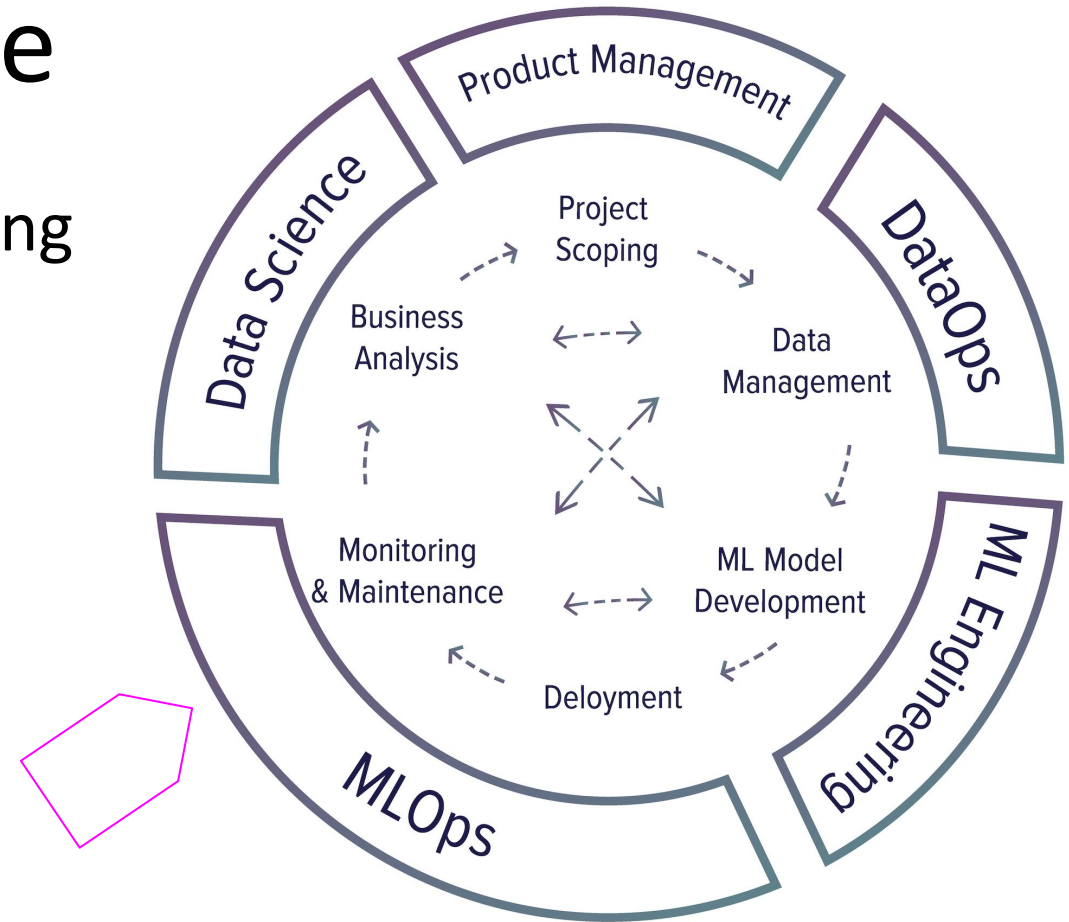
# Deployment

- Deploying and serving
- Release strategies
- Online model evaluation



# Monitoring & maintenance

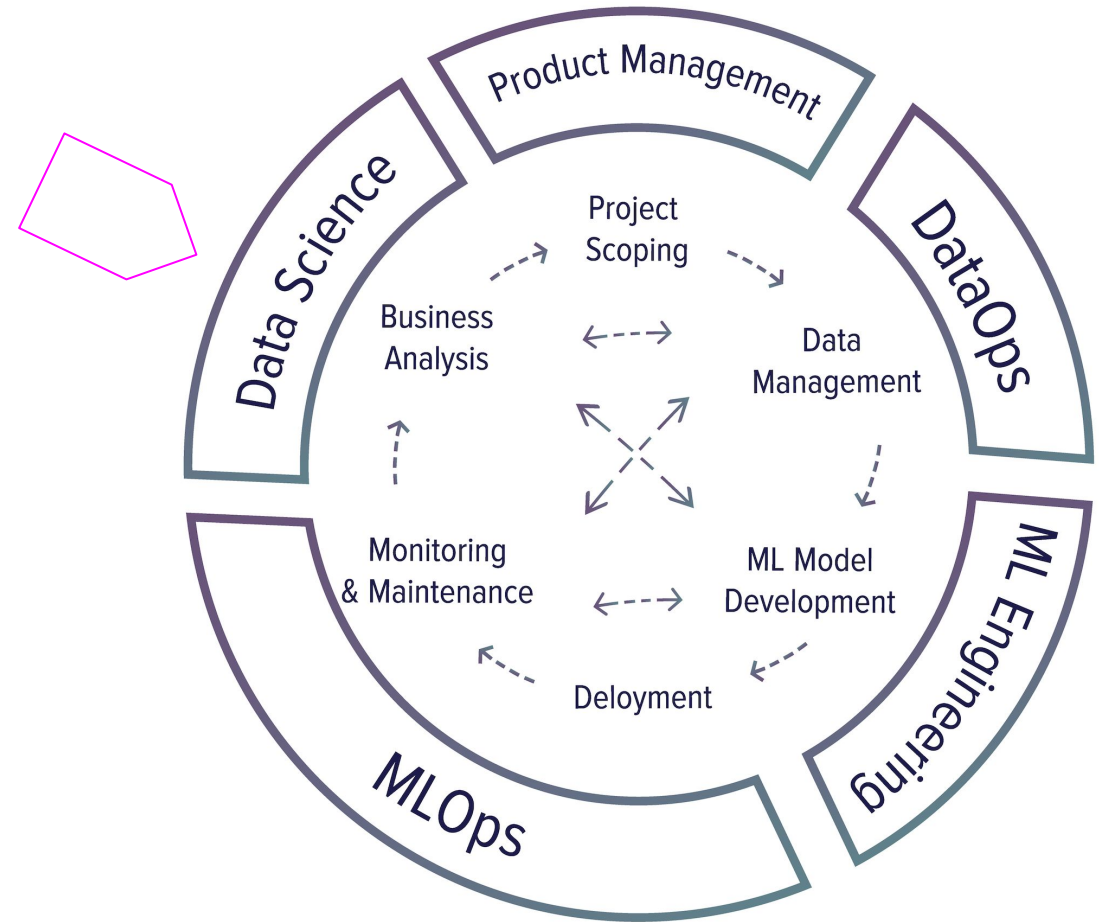
- Model performance & data monitoring
- Model retraining
- Model updates





# Business analysis

- User experience
- Tying model performance to business performance

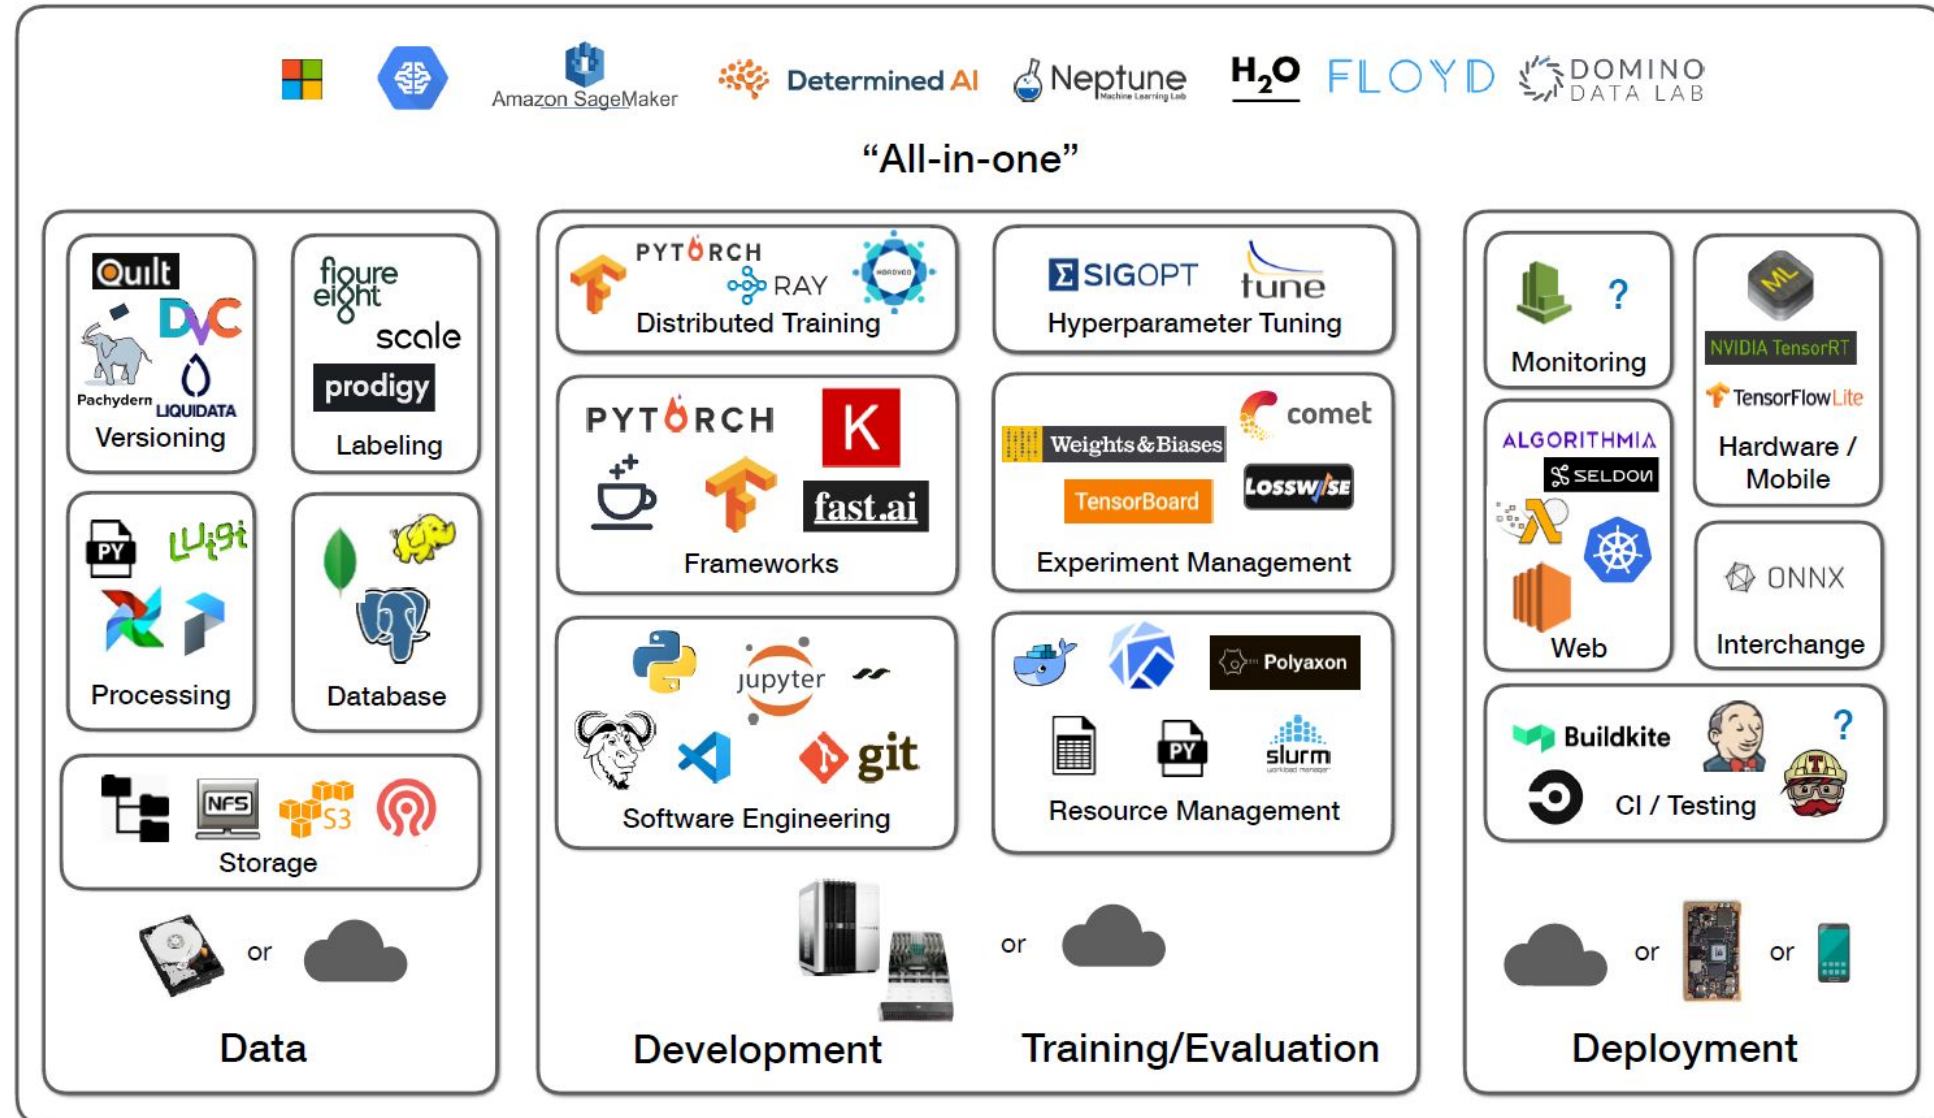




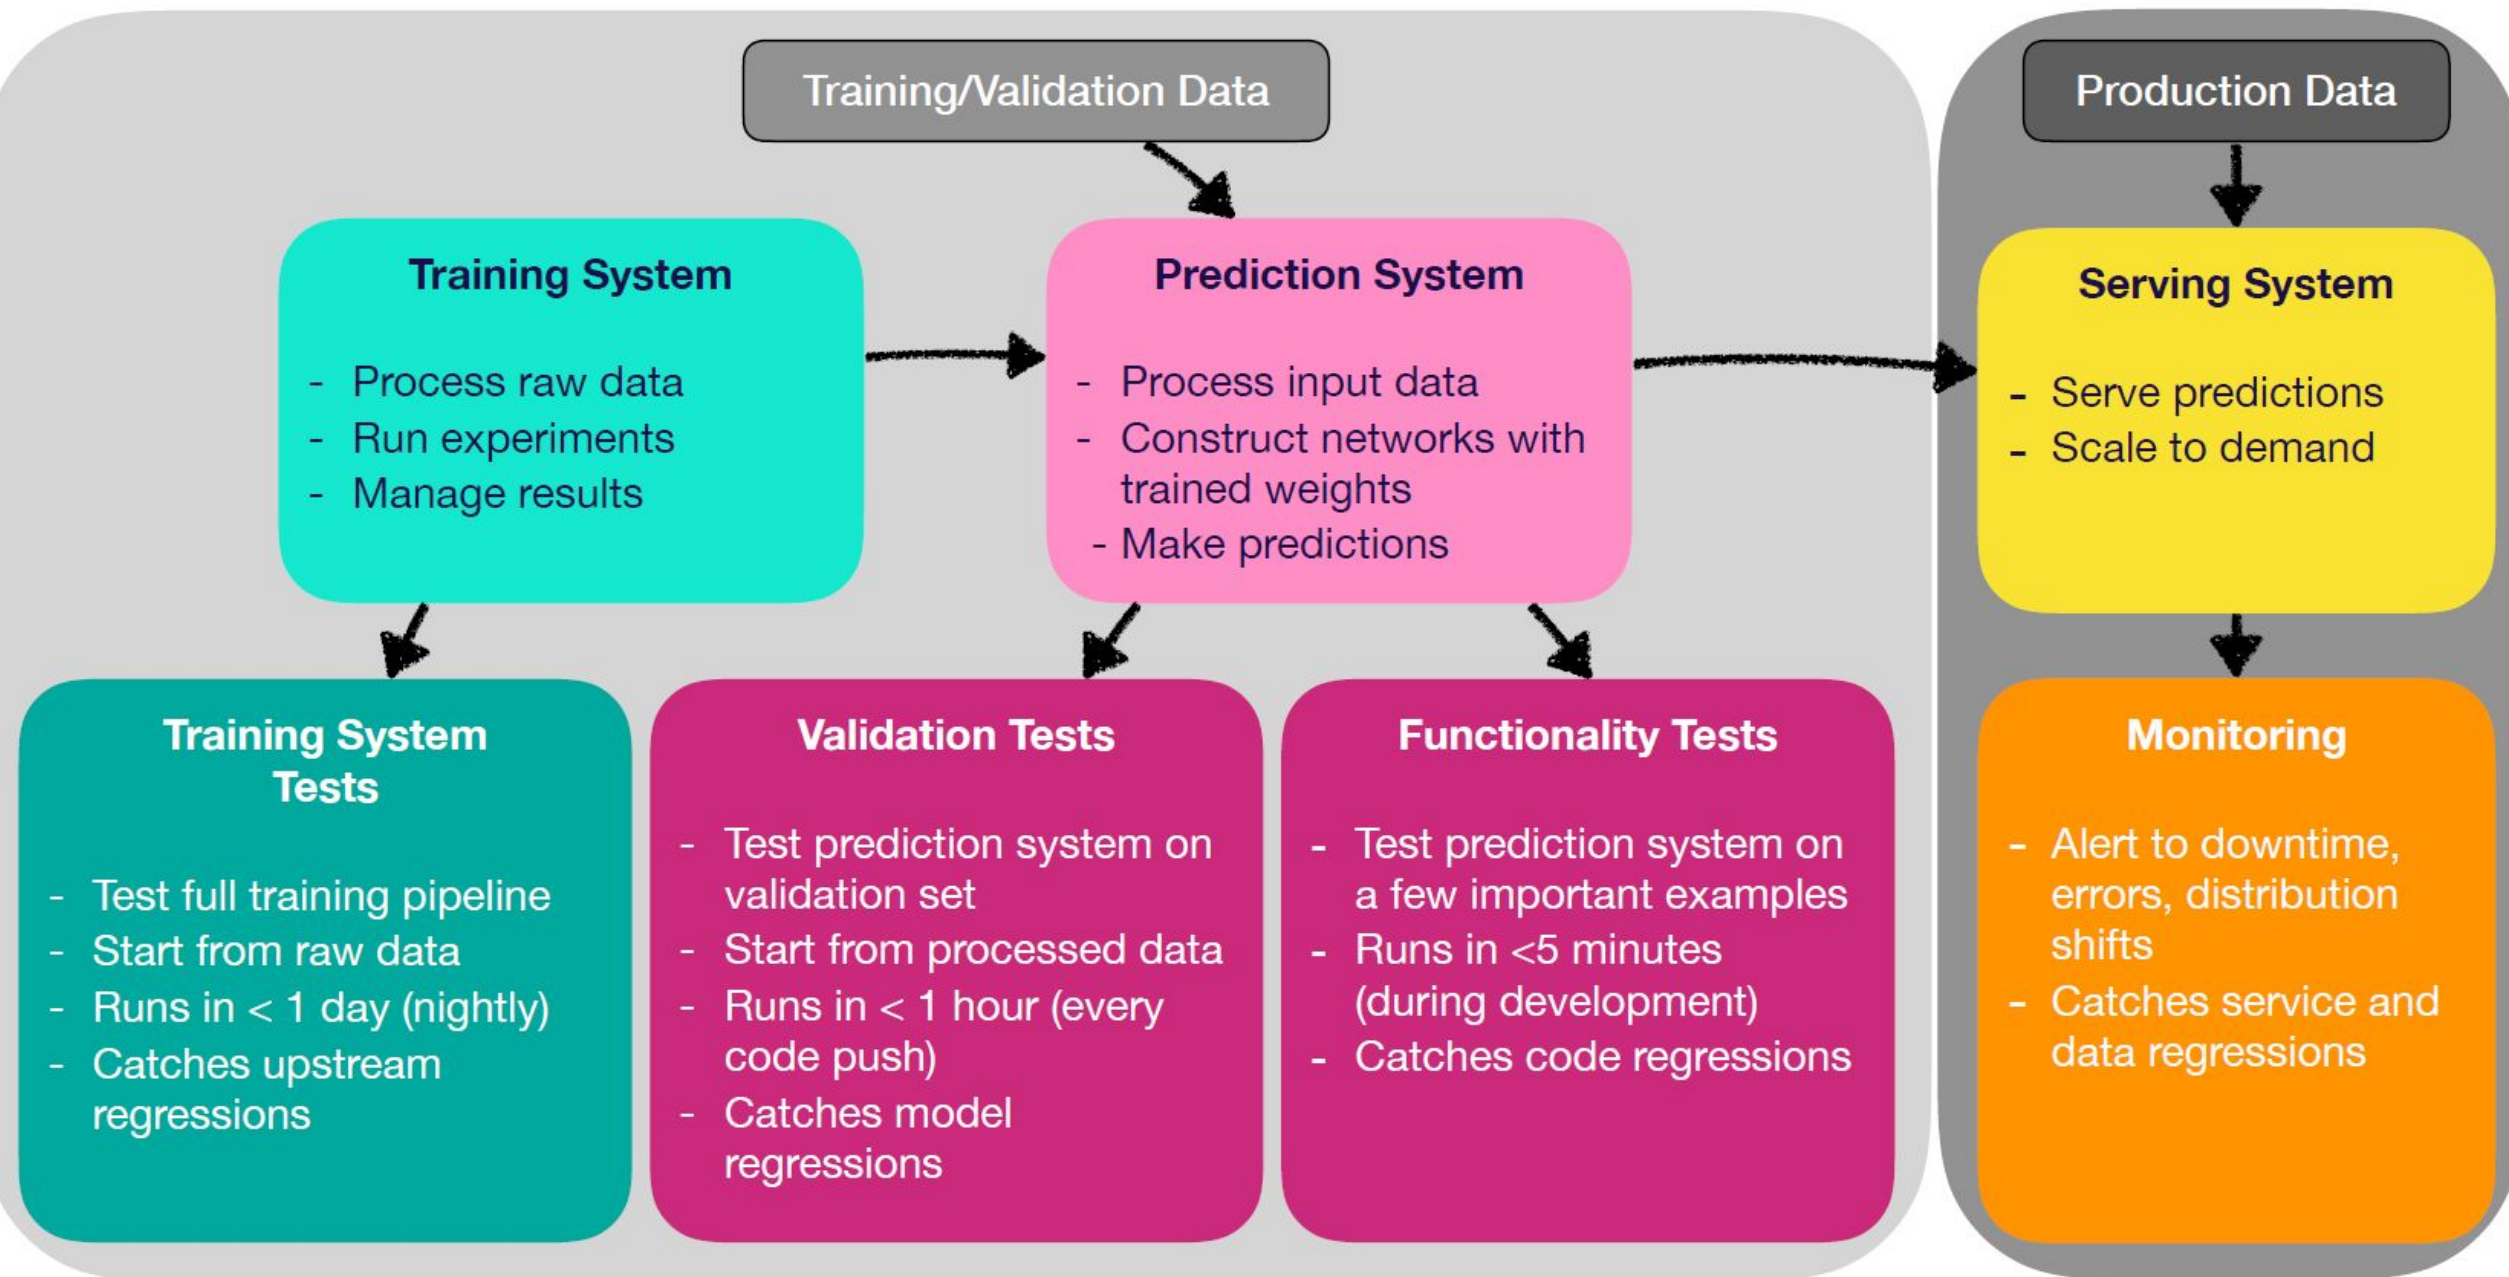
# Research Vs Production

		Production
Objectives		Different stakeholders have different objectives
Computational priority		Fast inference, low latency
Data		Constantly shifting
Fairness		Important
Interpretability		Important

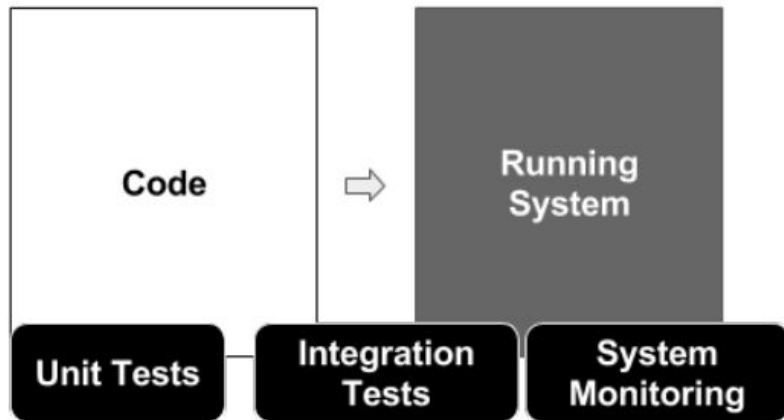
# ML Ops Stack



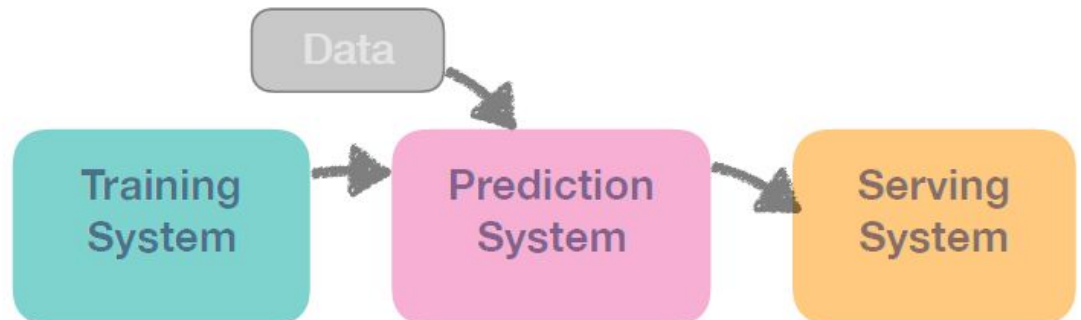
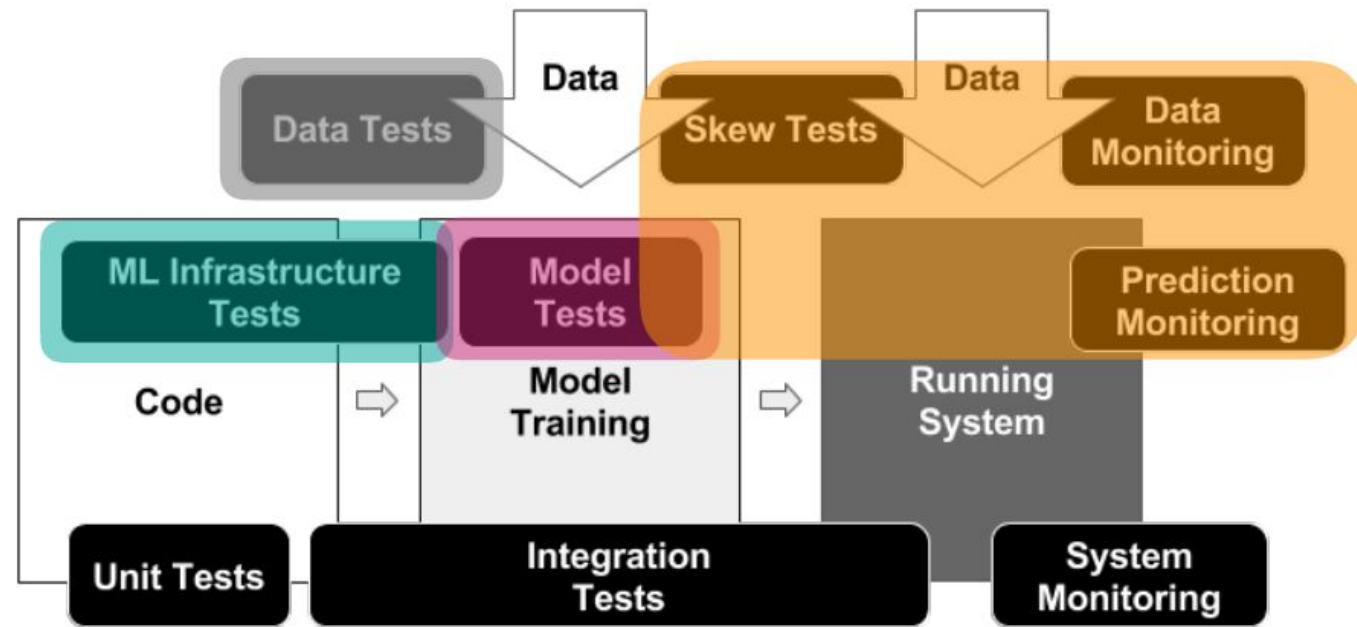
# Testing ML Pipelines



## Traditional Software



## Machine Learning Software





- |   |   |
|---|---|
| 1 | Feature expectations are captured in a schema.      |
| 2 | All features are beneficial.                        |
| 3 | No feature's cost is too much.                      |
| 4 | Features adhere to meta-level requirements.         |
| 5 | The data pipeline has appropriate privacy controls. |
| 6 | New features can be added quickly.                  |
| 7 | All input feature code is tested.                   |

### Data Tests

- |   |   |
|---|---|
| 1 | Model specs are reviewed and submitted.               |
| 2 | Offline and online metrics correlate.                 |
| 3 | All hyperparameters have been tuned.                  |
| 4 | The impact of model staleness is known.               |
| 5 | A simpler model is not better.                        |
| 6 | Model quality is sufficient on important data slices. |
| 7 | The model is tested for considerations of inclusion.  |

### Model Tests

- |   |  |
|---|--|
| 1 | Training is reproducible.                  |
| 2 | Model specs are unit tested.               |
| 3 | The ML pipeline is Integration tested.     |
| 4 | Model quality is validated before serving. |
| 5 | The model is debuggable.                   |
| 6 | Models are canaried before serving.        |
| 7 | Serving models can be rolled back.         |

### ML Infrastructure Tests

- |   |  |
|---|--|
| 1 | Dependency changes result in notification. |
| 2 | Data invariants hold for inputs.           |
| 3 | Training and serving are not skewed.       |
| 4 | Models are not too stale.                  |
| 5 | Models are numerically stable.             |
| 6 | Computing performance has not regressed.   |
| 7 | Prediction quality has not regressed.      |

### Monitoring Tests

# Further Reading

Best Practices for ML Engineering:

<https://developers.google.com/machine-learning/guides/rules-of-ml>

# References

- <https://stanford-cs329s.github.io/index.html>
- <https://fall2019.fullstackdeeplearning.com/>